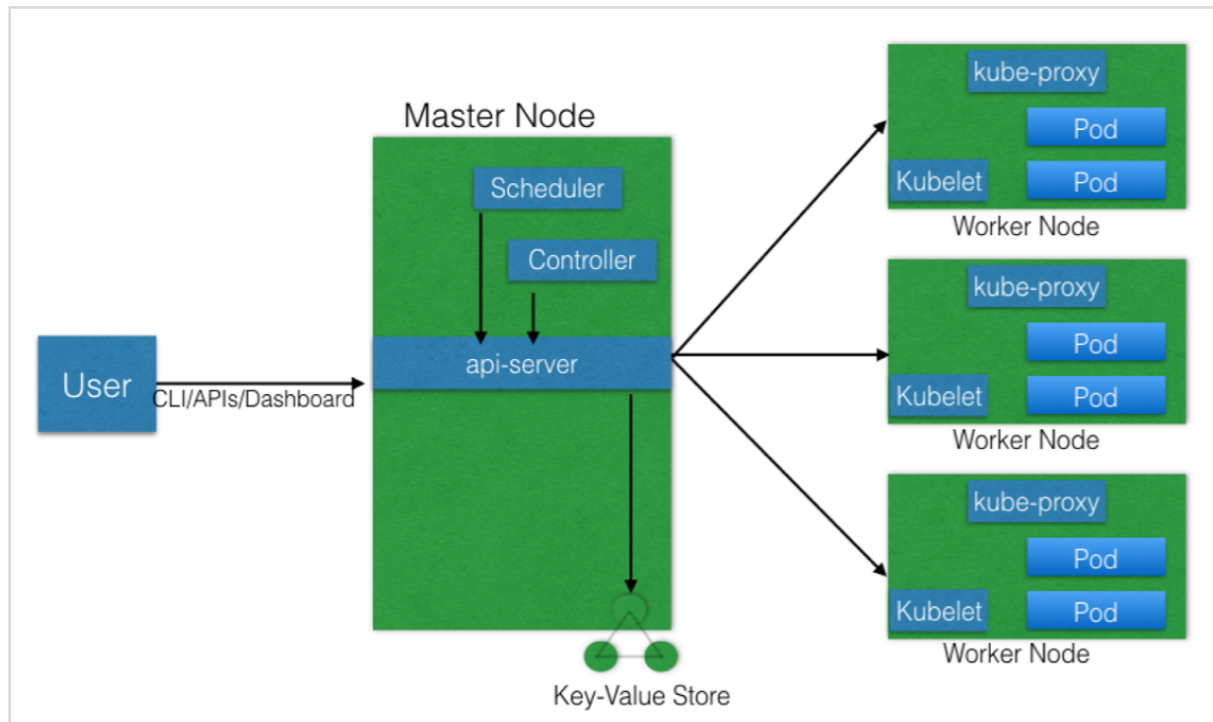


Chapter 4 Kubernetes Architecture

Master node

Worker Node

Networking Challenges



Master node

- Provides a running environment for the control plane
 - Manages the state of a Kubernetes cluster
 - Brain behind all operations inside the cluster
- User interacts with master node
- Replicas exist in the cluster with sync'd control plane components
 - Minimize downtime

API Server (kube-apiserver)

- Intercepts, validates, and processes RESTful calls
- talks to etcd data store (fetches and changes state)
- Custom API servers
 - Primary API server becomes a proxy to all secondary custom ones
 - Routes the calls to them based on custom defined rules

Scheduler (kube-scheduler)

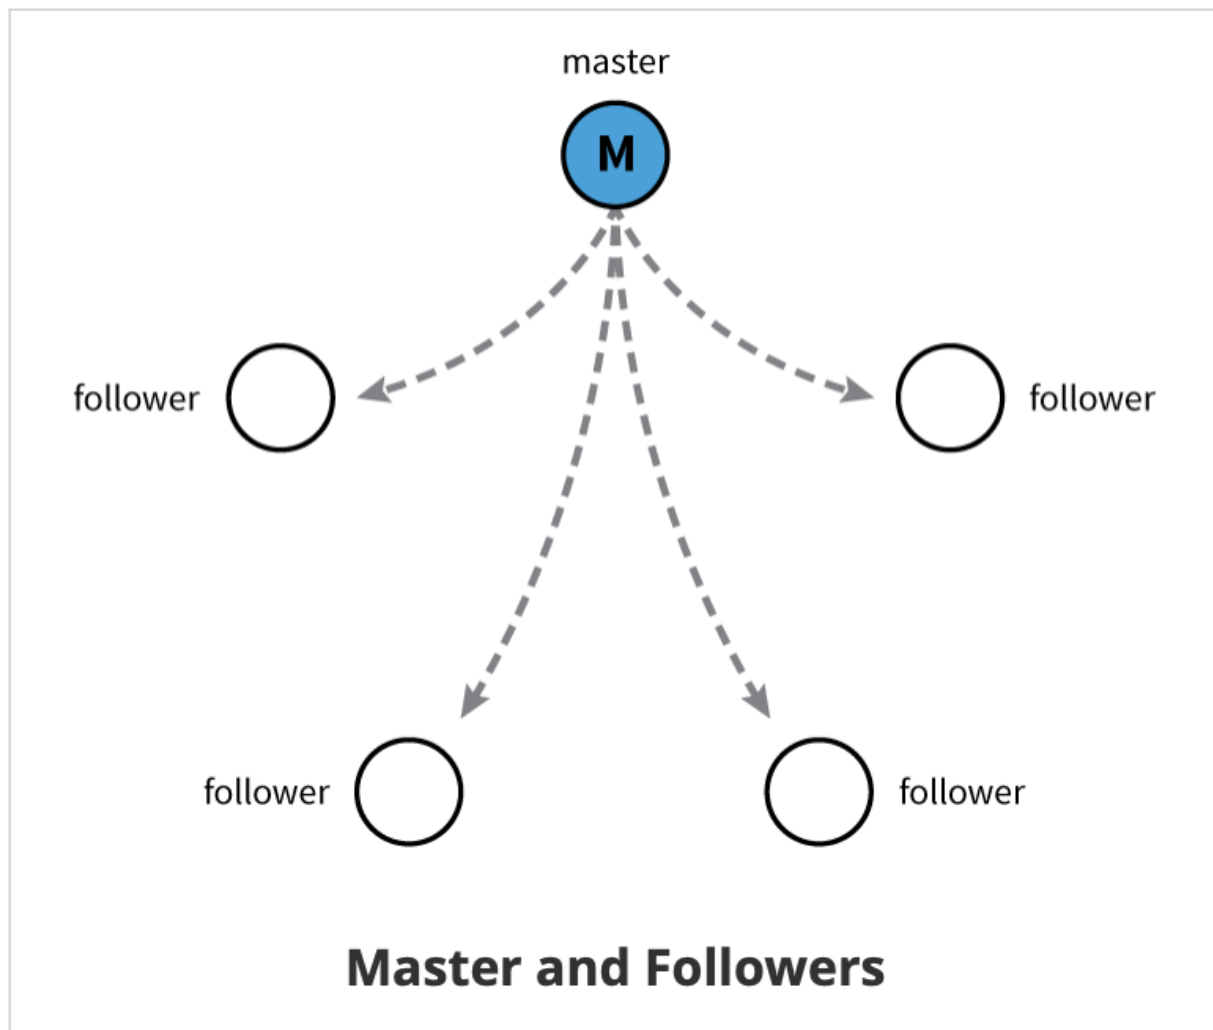
- Assigns new objects to nodes
- Talks to API server to learn about the new state in etcd

Controller managers

- Regulates the state of cluster
- Kube-controller-manager checks nodes'
 - Availability
 - Pod count
 - Endpoints
 - Service accounts
 - API access tokens
- Cloud-controller-manager
 - Interact with cloud provider when nodes become unavailable
 - Manage storage volumes provided by cloud service
 - Manage LB and routing

Etcd: Distributed Key-Value ****State Store****

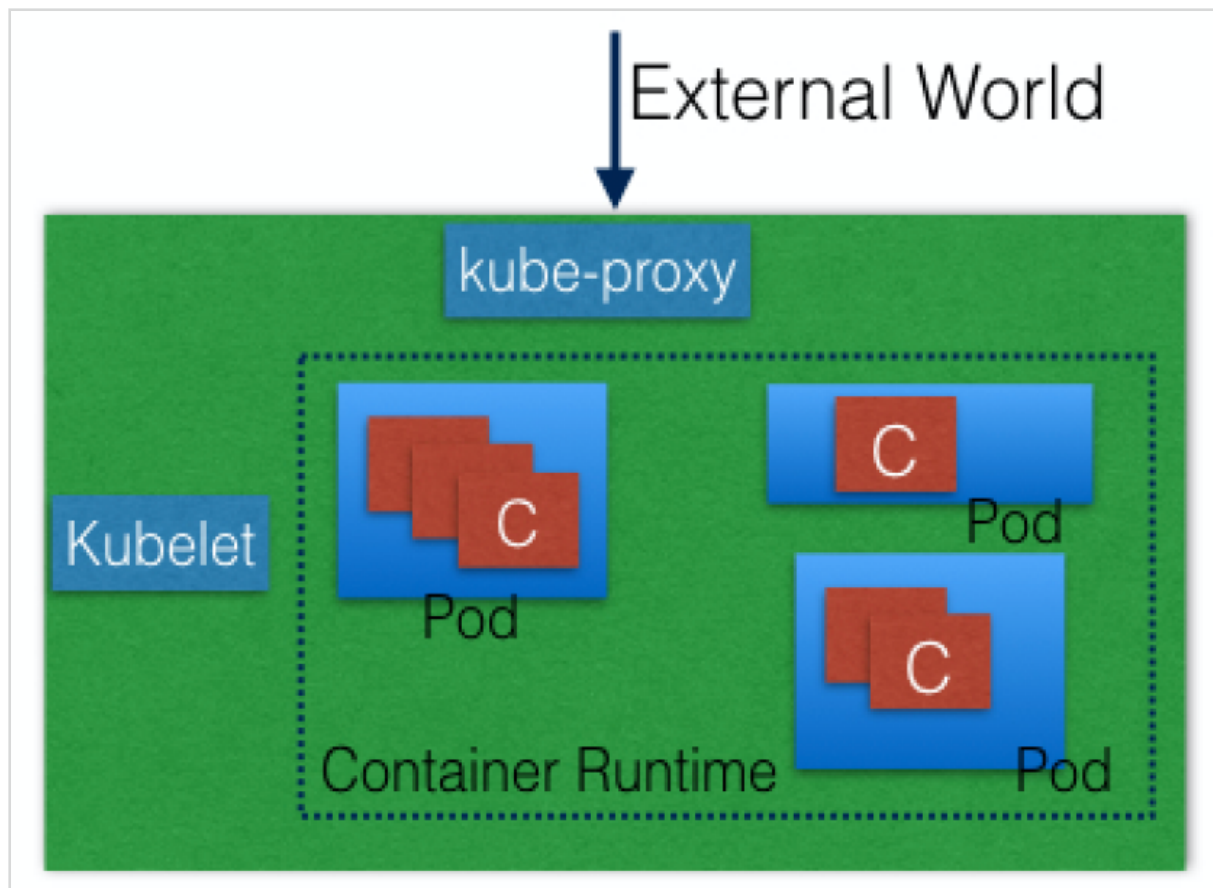
- Holds cluster state related data, not client workload data
- Stacked: on master node
- External: on dedicated host
- New data is appended, never replaced
- Obsolete data is compacted periodically
- Very important to replicate data stores in HA mode
- Raft Consensus Algorithm



- Failure protection
- Any one node as master, others as followers
- Written in Go
- Stores subnets, ConfigMaps, Secrets, etc.

Worker Node

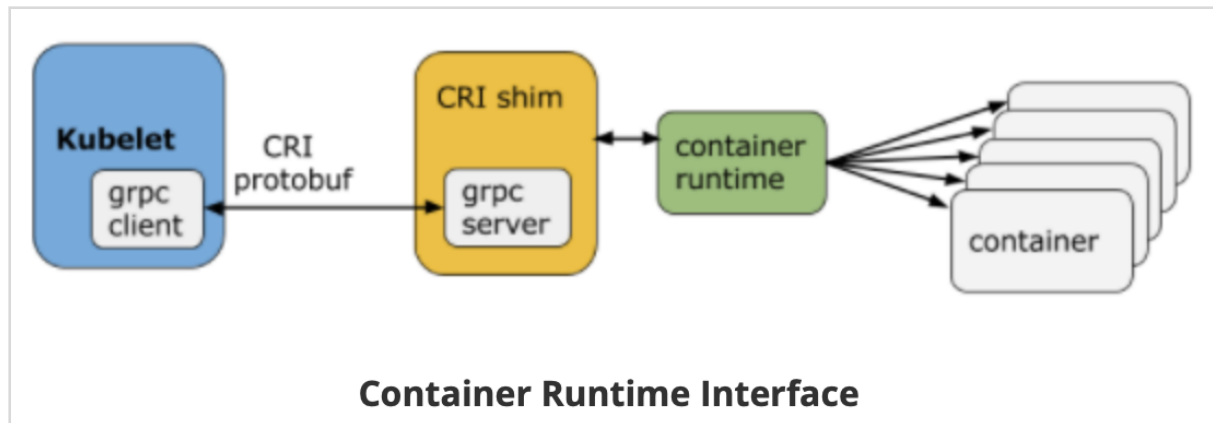
[back to top](#)



- Environment for Client applications
- Apps encapsulated in pods
 - Controlled by the cluster control plane agents (on MN)
- Can talk to each other + outside world
- Pod: smallest scheduling unit

Kubelet


- Agent that communicates with the control plane components
- Receives pod definitions
- Interacts with the container runtime (CR) to run containers associated with the Pod
- Monitors health of pods' containers
- Connects to CR using Container Runtime Interface (CRI)



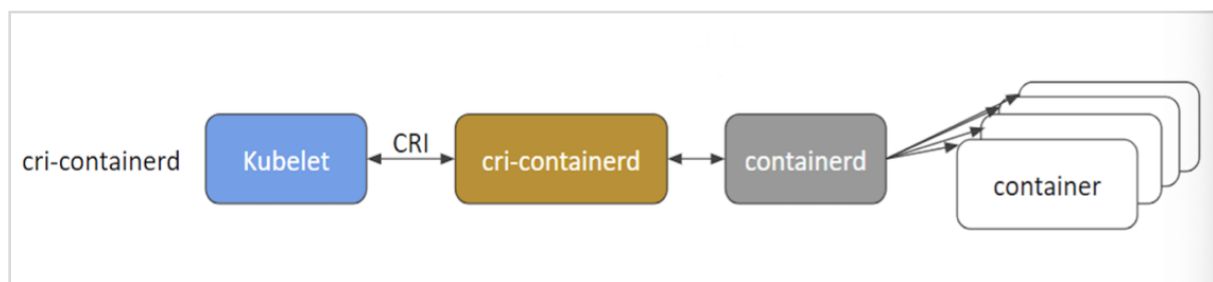
- kubelet acts as grpc client
- connects to CRI shim activating as grpc server
- CRI implements two services:
 - ImageService: image-related operations
 - RuntimeService: pod and container-related operations

CRI Shims

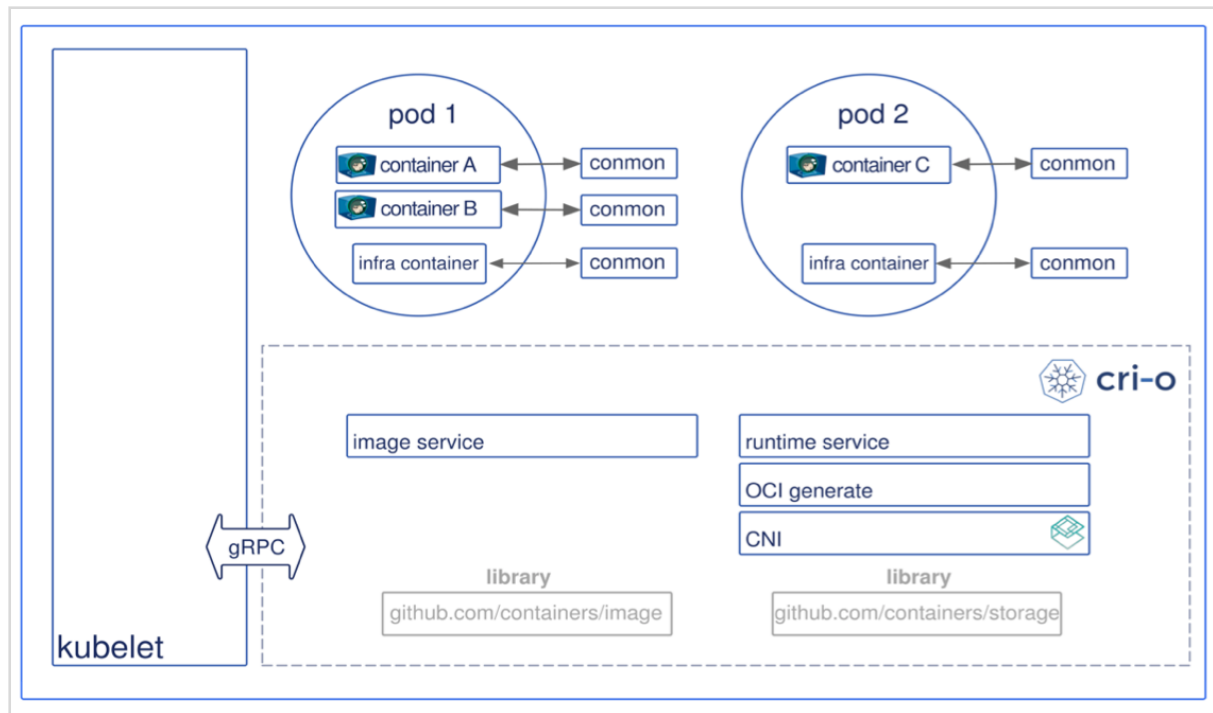
- Dockershim



- containers to create and manage containers
- Cri-containerd



- Directly use Docker's smaller offspring containerd
- CRI-O



- enables using any Open Container Initiative (OCI) compatible runtimes with Kubernetes

Kube-Proxy

- Network agent on each node
- Dynamic updates and maintenance of networking rules
- Forwards connection requests to pods

Add-ons

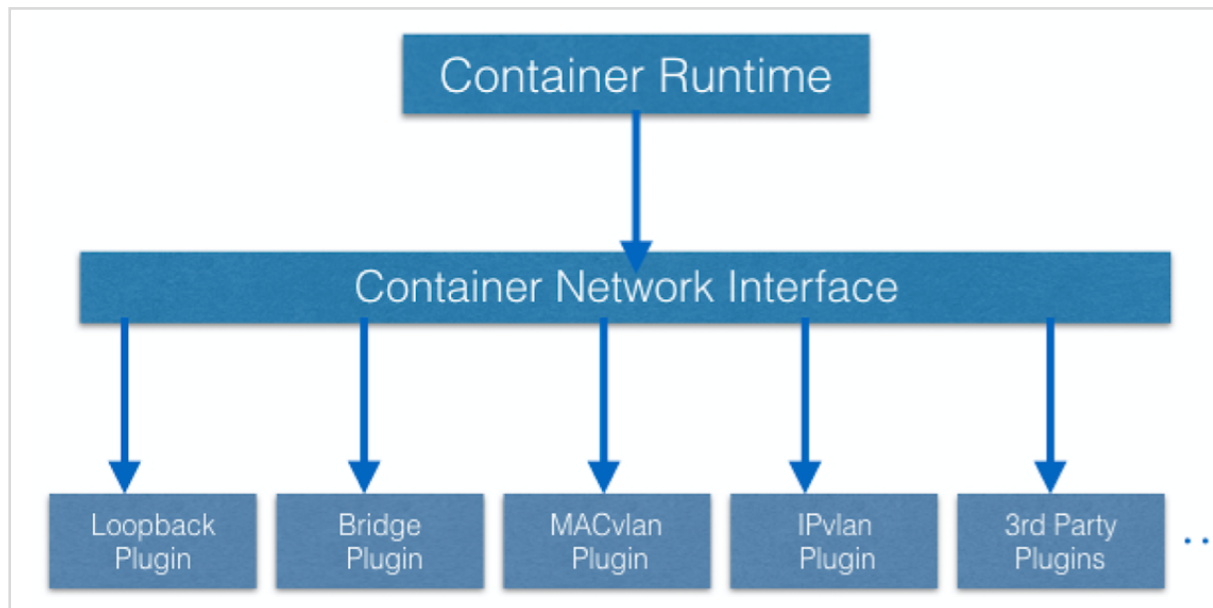
- 3rd party pods and services
- Cluster **DNS**, Dashboard, Monitoring, Logging

Networking Challenges

Back to top

- Container-to-Container inside Pods
 - Network Namespace
 - Created when a pod is started by container runtime
 - Shared across containers on the pod
 - Talk to each other via localhost
- Pod-to-Pod Communication Across Nodes
 - IP-per-pod: similar to VMs on a network
 - Containers inside pods coordinate port assignments
 - Container Network Interface (CNI)

- set of specification and libraries
- allows plugins to configure the networking for containers
- 3rd-party Software Defined Networking (SDN) solutions



- CR offloads the IP assignment to CNI
- CNI connects to the plugins and get the IP address
- CNI forwards the IPs to CR
- Pod-to-External World communication
 - Services
 - Enable external accessibility
 - Encapsulate networking rules definition on cluster nodes
 - Kube-proxy
 - Exposes services to the external world
 - Applications accessible by outside world over a virtual IP